

Resolution-Agnostic Learning: Neural Operators That Scale

Junbin Gao

The University of Sydney Business School

Sydney AI Meetup, 24 November 2025

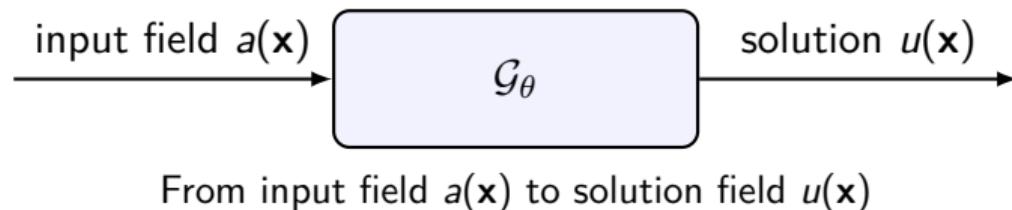
Outline

- 1 Motivation
- 2 Neural Operator Construction
- 3 Operators: Overall Architectures
- 4 Functional Tensor Decomposition
- 5 Time-aware Neural Operators
- 6 Neural Operators for SDE/SPDEs

Motivation

Why neural operators?

- Expensive simulators (PDEs, multi-physics) \Rightarrow need fast surrogates.
- Standard nets learn point-to-point maps $\mathbb{R}^d \rightarrow \mathbb{R}^{out}$; **neural operators** learn maps *between functions*.



- **Resolution-agnostic:** train on coarse meshes; infer on fine without retraining.
- **Generalize** across parameters, geometries, forcing.
- In general, functions a and u may have different domains. Because of this we can use notation $u(\mathbf{y})$

Neural Operators (NOs)

- Neural operators are a powerful class of models designed to learn mappings between infinite-dimensional function spaces, making them well-suited for functional regression problems and solving parametric partial differential equations (PDEs) in a resolution-agnostic way.
- Typically, given two Banach spaces \mathcal{A}, \mathcal{U} , the “input-output” relationship is an operator $F: \mathcal{A} \rightarrow \mathcal{U}$ acting between these two function spaces.
- A neural operator F_θ , with learnable parameters θ , serve as a continuous approximation of F . It is independent of any particular input discretization, allowing a single model to generalize consistently across different meshes.
- Specifically, the architecture F_θ has two key properties: **1) universality** and **2) discretization-invariant**.

Operator intuition: functional analysis

- Functional view of vectors: A d -dimensional vector $\mathbf{x} = (x_1, x_2, \dots, x_d) \in \mathbb{R}^d$ can be regarded as a function

$$x : \mathcal{D} = \{1, 2, \dots, d\} \rightarrow \mathbb{R}, \quad x(i) = x_i.$$

- Vector view of functions: A function a defined on domain, e.g., $\mathcal{D} = [0, 1]$, is a vector.

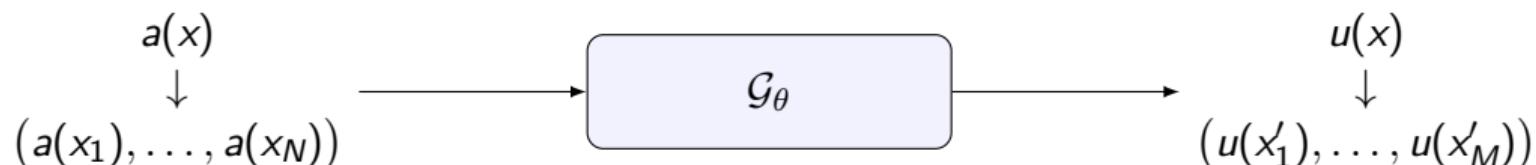
$$a : \mathcal{D} = [0, 1] \rightarrow \mathbb{R}, \quad \forall x \in \mathcal{D}, y = a(x) \in \mathbb{R}$$

by thinking of $\mathcal{D} = \{x_1, x_2, x_3, \dots\}$, the function can be a (long) vector
 $a = (a(x_1), a(x_2), a(x_3), \dots)$. a.k.a function value table.

- Education: A function is a with its domain \mathcal{D} , not $a(x)$ where $a(x)$ is the calculation rule for the function or mapping rule.
- An **operator** maps one function to another: $\mathcal{G} : a \mapsto u$, which means mapping a (long) vector to another (long) vector. Exactly what a neural network does!

Samples of functions

- In application, a function is normally presented as a finite number of samples in terms of $\{(\mathbf{x}_i, u_i)\}_{i=1}^N$ where $\mathbf{x}_i \in \mathcal{D}$ and $u_i \in \mathbb{R}$. Or when \mathbf{x}_i is clear (e.g. uniformed discrete points on \mathcal{D}), then the function is presented by N -dimension vector $(u_1, u_2, \dots, u_N) \in \mathbb{R}^N$.
- **Neural operators** learn this map *without fixing N or M*: same parameters work across meshes/resolutions.



Same θ works for different N, M (resolutions).

- How to design such “operators”?

From Neural Networks to Neural Operators

- Start: standard NN is **vector**→**vector**

$$f_\theta : \mathbb{R}^N \rightarrow \mathbb{R}^M, \quad \mathbf{y} = f_\theta(\mathbf{x})$$

Make “functions” into **vectors** by sampling on grids:

$$a(\mathbf{x}) \Rightarrow \mathbf{a} \in \mathbb{R}^N, \quad u(\mathbf{x}) \Rightarrow \mathbf{u} \in \mathbb{R}^M.$$

- Goal: learn a *single* f_θ that works across resolutions (different N, M , different $\{x_1, \dots, x_N\}$ and $\{x'_1, \dots, x'_M\}\} \Rightarrow$ **neural operator**.

- Key ideas

- Parameter sharing across grids: same θ for coarse/fine meshes.
- Coordinate-aware features (positional encodings, trunk nets).
- Physics-friendly priors (e.g., Fourier layers).

Neural Operator Construction

Examples of NO Architectures

- Deep Operator Network (DeepONet), learning nonlinear operators by combining a “branch” network that encodes sampled values of the input function with a “trunk” network that encodes evaluation locations to produce the operator’s output.
- Graph Neural Operator (GNO), generalizing graph neural networks to learn “linear” mappings between functions, aka *integral operators*, by iteratively applying message-passing layers whose weights are shared across nodes (location of input or output function values) and conditioned on positional or edge features, thus approximating nonlinear operators
- Fourier Neural Operator (FNO), learning nonlinear operators on continuous domains by repeatedly lifting input fields into spectral space, applying learned complex-valued multipliers to selected Fourier modes, and transforming back to physical space, enabling mesh-independent prediction of function-to-function mappings.

Two Practical Constructions

DeepONet (basis expansion)

$$u(\mathbf{x}) \approx \sum_{r=1}^R \underbrace{b_r(\mathbf{a})}_{\text{branch}} \underbrace{t_r(\mathbf{x})}_{\text{trunk}}$$

Works on arbitrary point sets; resolution-agnostic via \mathbf{x} -query.

FNO (spectral convolution)

$$\hat{u}'(\mathbf{k}) = W(\mathbf{k}) \hat{u}(\mathbf{k}) \text{ on low modes}, \quad u' = \mathcal{F}^{-1}[\hat{u}'] + \text{skip}$$

Grid-agnostic via FFT basis; strong bias for PDE dynamics.

Outcome: a learned map $\mathcal{G}_\theta : \mathbf{a} \mapsto u$ that generalizes across meshes, domains, and parameters.

DeepONet Branch Design (Resolution-Invariant Encoders)

Goal: map sampled input function a on *any* grid $\{(x_j, a_j)\}_{j=1}^N$ to branch features $b(a) \in \mathbb{R}^R$ that are *independent of N and spacing*.

① DeepSets / Sample–then–Pool

$$b(a) = \rho \left(\sum_{j=1}^N \phi(a_j, x_j) w_j \right)$$

Permutation-invariant; w_j can be quadrature weights. ϕ, ρ are small MLPs.

② Attention Pooling (Set Transformer)

$$\alpha_j = \text{softmax}(g(a_j, x_j)), \quad b(a) = \sum_{j=1}^N \alpha_j h(a_j, x_j)$$

Learns where to “look”; handles nonuniform sampling.

DeepONet Branch Design (Resolution-Invariant Encoders)

① Projection onto (Learned or Fixed) Bases

$$b_r(a) = \sum_{j=1}^N a_j \psi_r(x_j) w_j, \quad \psi_r \in \{\text{Fourier/SIREN/poly or learned}\}$$

Acts like resolution-agnostic features via quadrature.

② Graph Encoder (Message Passing)

Build k -NN graph on $\{x_j\}$; run L message-passing layers, then global pool:

$$z = \text{Pool}(\text{GNN}_\theta(\{(x_j, a_j)\})), \quad b(a) = Wz$$

Robust to irregular meshes and varying density.

③ Spectral Truncation Encoder (FNO-style)

FFT on the input grid, keep low modes $|k| \leq K$ (fixed K), iFFT or linear map:

$$b(a) = \text{vec}(\hat{a}(k))_{|k| \leq K}$$

Fixed number of retained modes \Rightarrow fixed branch size across resolutions.

DeepONet Branch Design (Resolution-Invariant Encoders)

- ① **CNN with Adaptive Pooling (for gridded inputs)** Apply small CNN/ResNet, then *adaptive global average pooling* to a fixed size:

$$b(a) = \text{AGAP}(\text{CNN}(a))$$

Practical when inputs live on images/structured grids.

- ② **Trunk (coordinate net):** $t(x)$ can be a standard MLP/SIREN with positional encodings.

DeepONet output: $u(x) \approx \sum_{r=1}^R b_r(a) t_r(x)$.

From Linear Layers to Linear Operators (Kernel View)

Finite-dimensional (NN) view

$$\mathbf{u} = W\mathbf{x} \quad (\text{ignore bias})$$

where $W \in \mathbb{R}^{M \times N}$ is a *linear layer*.

Function-space (operator) view

$$u(x) := (\mathcal{K}a)(x) = \int_{\mathcal{D}} k(x, \xi) a(\xi) d\xi$$

A *linear operator* applied to a function a ; $k(x, \xi)$ is the **kernel** (simulating W).

Discretize functions as long vectors

$$a(\xi_\ell) \Rightarrow \mathbf{a} \in \mathbb{R}^N, \quad u(x_j) \Rightarrow \mathbf{u} \in \mathbb{R}^M.$$

Integral \Rightarrow matrix multiply

$$u(x_j) \approx \sum_{\ell=1}^N k(x_j, \xi_\ell) a(\xi_\ell) w_\ell$$

Let $K_{j\ell} = k(x_j, \xi_\ell)$ and $W_{\ell\ell} = w_\ell$ (quadrature weights):

$$\underbrace{\mathbf{u}}_{\mathbb{R}^M} \approx \underbrace{K}_{\mathbb{R}^{M \times N}} \underbrace{W}_{\mathbb{R}^{N \times N}} \underbrace{\mathbf{a}}_{\mathbb{R}^N} \iff \mathbf{u} = \tilde{W}\mathbf{a}.$$

Takeaway: a linear operator on functions becomes that a *kernel matrix* times a *function vector* after sampling.

One way to implement this kernel is to use an NN to compute $k(x, \xi)$ for a pair of x and ξ .

Special case: Convolution

- Choose the kernel as $k(x, \xi) = h(x - \xi)$, then

$$u(x) = (\mathcal{K}a)(x) = \int h(x - \xi) a(\xi) d\xi = (h * a)(x) \Rightarrow \text{Resulting Toeplitz/Circulant } \tilde{W}.$$

$$T = \begin{pmatrix} 3 & 2 & 5 & 7 & 4 \\ 6 & 3 & 2 & 5 & 7 \\ 8 & 6 & 3 & 2 & 5 \\ 9 & 8 & 6 & 3 & 2 \\ 1 & 9 & 8 & 6 & 3 \end{pmatrix} \quad \text{and} \quad C = \text{circ}(1, 2, 3, 4, 5) = \begin{pmatrix} 1 & 2 & 3 & 4 & 5 \\ 5 & 1 & 2 & 3 & 4 \\ 4 & 5 & 1 & 2 & 3 \\ 3 & 4 & 5 & 1 & 2 \\ 2 & 3 & 4 & 5 & 1 \end{pmatrix}$$

- As $u(x) = (h * a)(x) \Leftrightarrow \hat{u}(k) = \hat{h}(k) \cdot \hat{a}(k)$, This choice is equivalent to design Fourier Operator layers.

From Kernel Operators to FNO (Fourier View)

Linear operator as integral:

$$u(x) = (\mathcal{K}a)(x) = \int_{\mathcal{D}} k(x, \xi) a(\xi) d\xi$$

If translation-invariant $k(x, \xi) = \kappa(x - \xi)$ (convolution):

$$\widehat{(\mathcal{K}a)}(k) = \widehat{\kappa}(k) \widehat{a}(k)$$

Discretize & learn in frequency:

$$\hat{u}'(k) = W_\theta(k) \hat{u}(k) \quad (\text{per low-frequency mode})$$

Keep only $|k| \leq K$ (spectral truncation) for efficiency/inductive bias.

One FNO layer (practical form):

$$u^{(\ell+1)} = \sigma \left(W_0 u^\ell + \mathcal{F}^{-1} [R_\theta(k) (\mathcal{F} u^\ell)(k)]_{|k| \leq K} \right)$$

- $R_\theta(k)$: learnable complex weights per mode (the kernel in Fourier space).
- W_0 : pointwise mixing (skip/linear).
- σ : nonlinearity; residual stacking builds *nonlinear* operator.

Why this yields a neural operator

- Parameters live on modes \Rightarrow *resolution-agnostic*.
- Convolution theorem encodes PDE-friendly inductive bias.

From Kernel Integrals to Graph Neural Operators (GNO)

- **Function → Graphs: bipartite first, then primal**

$$u(x) = \int_{\Omega} k(x, \xi) a(\xi) d\xi$$

- **Input discretization**

$$\{x_i\}_{i=1}^M \text{ for } u, \quad \{\xi_j\}_{j=1}^N \text{ for } a$$

- Bipartite graph $\mathcal{G}_{a \rightarrow u}$ between u -nodes x_i and a -nodes ξ_j .
- Edges: $j \in \mathcal{N}_a(i)$ by k -NN/radius in $\|x_i - \xi_j\|$.
- Edge feats e_{ij} : $(x_i - \xi_j, \|x_i - \xi_j\|, \text{BC/geom})$, weights w_{ij} (quadrature/normalization).

- **After the first layer**

- Build a *primal u-graph* $\mathcal{G}_{u \leftrightarrow u}$ on $\{x_i\}$ (k NN/radius).
- Optionally keep (or refresh) a cross-graph $\mathcal{G}_{a \rightarrow u}$ for re-conditioning.

- **Resolution-agnostic:** same kernels on new meshes; only neighborhoods change.

From Kernel Integrals to Graph Neural Operators (GNO)

- **Layer 1: bipartite (encode $a \rightarrow u$)** (passing $\{a(x_i)\}$ to nodes $\{\xi_j\}$)

$$u_i^{(1)} = \sigma \left(b_u + \sum_{j \in \mathcal{N}_a(i)} k_\theta^{(0)}(x_i, \xi_j, e_{ij}) a_j w_{ij} \right)$$

- **Layers $\ell \geq 1$: primal u -graph propagation** (all nodes $\{x_i, \xi_j\}$)

$$u_i^{(\ell+1)} = \sigma \left(W^{(\ell)} u_i^{(\ell)} + \sum_{p \in \mathcal{N}_u(i)} g_\theta^{(\ell)}(x_i, x_p, e_{ip}) u_p^{(\ell)} \right)$$

- **Design choices**

- *Hybrid (default)*: 1 bipartite layer, then L primal layers.
- *Alternating*: $(a \rightarrow u)$ then $(u \leftrightarrow u)$, repeat.
- *Shared grid*: if $x_i \equiv \xi_i$, the first layer reduces to a standard GNN conv.
- *Attention form*: normalize $k_\theta^{(0)}$ or $g_\theta^{(\ell)}$ with softmax over neighbors.

From Linear GNO to Nonlinear Kernel Neural Operators

- **Operator viewpoint**

$$(\mathcal{K}[a, u^{(\ell)}] v)(x) = \int_{\Omega} k_{\theta}(x, \xi; a(\xi), u^{(\ell)}(x), u^{(\ell)}(\xi)) v(\xi) d\xi$$

- If k_{θ} does *not* depend on a, u^{ℓ} \Rightarrow linear kernel.
- If k_{θ} depends on (a, u^{ℓ}) \Rightarrow **nonlinear, input-/state-dependent** operator.

- **Discretization (bipartite then primal)**

$$u_i^{(1)} = \sigma \left(\sum_{j \in \mathcal{N}_a(i)} k_{\theta}^{(0)}(x_i, \xi_j; a_j) a_j w_{ij} \right)$$

$$u_i^{(\ell+1)} = \sigma \left(\mathcal{W}^{(\ell)} u_i^{(\ell)} + \sum_{p \in \mathcal{N}_u(i)} k_{\theta}^{(\ell)}(x_i, x_p; a_p, u_i^{(\ell)}, u_p^{(\ell)}) u_p^{(\ell)} w_{ip} \right)$$

Recovery: remove dependence on (a, u^{ℓ}) to get the linear GNO layer.

From Linear GNO to Nonlinear Kernel Neural Operators

- Three useful nonlinear kernel designs

- ① Attention (*nonlocal*) kernel

$$\begin{aligned} k_\theta &= \text{softmax}_{p \in \mathcal{N}_u(i)}(q_\theta(x_i, u_i^{(\ell)})^\top k_\theta(x_p, u_p^{(\ell)})) \\ \Rightarrow u_i^{(\ell+1)} &= \sum_p k_\theta(\cdot) v_\theta(u_p^{(\ell)}) \end{aligned}$$

(cross-attn): replace u^ℓ at source with a .

- ② Dynamic (*space-variant*) convolution

$$k_\theta(x_i, x_p; \cdot) = \sum_{m=1}^M \beta_m^{(\ell)}(x_i, u_i^{(\ell)}) h_m(x_i - x_p)$$

Apply M fixed (FFT-able) kernels h_m then mix by β_m .

- ③ Neural RKHS kernel (*feature map*)

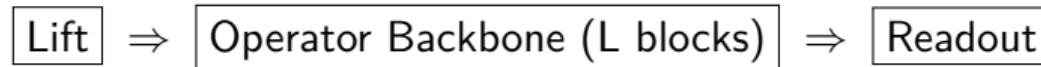
$$k_\theta(x_i, x_p; \cdot) = \phi_\theta(x_i, u_i^{(\ell)}, a_i)^\top \phi_\theta(x_p, u_p^{(\ell)}, a_p)$$

Positive definite by construction; control via $\|\phi_\theta\|$.

Operators: Overall Architectures

Neural Operators: Overall Architecture

- **Goal.** Learn an operator $\mathcal{G}_\theta : a \mapsto u$ that maps input function a on a domain \mathcal{D} to output function u , *independent of discretization*, pipelined in three stages:



- **Lift (local embedding, resolution-agnostic)**

$$u^{(0)}(x) = P(a(x), x, \text{PE}(x), \text{BC/geom}) \in \mathbb{R}^C$$

where P is a pointwise map (e.g. 1×1 conv/MLP). Include coordinates or positional encodings (PE) and boundary/geometry tags. Works on grids or point clouds.

- **Operator Backbone (nonlocal field mixing)**

$$u^{(\ell+1)} = \sigma \left(\mathcal{A}_\theta^{(\ell)} [u^{(\ell)}] + B^{(\ell)} u^{(\ell)} \right), \quad \ell = 0, \dots, L-1$$

with residual $B^{(\ell)}$ (pointwise).

Neural Operators: Overall Architecture

- The operator block $\mathcal{A}_\theta^{(\ell)}[\cdot]$ is *one of*:
 - **FNO (spectral)**: FFT \Rightarrow low-mode multipliers $\mathcal{W}^{(\ell)}(k) \Rightarrow$ IFFT.
 - **GNO (graph/point cloud)**: message passing on k -NN / radius graphs:

$$(\mathcal{A}_\theta^{(\ell)} u^{(\ell)})_i = \sum_{j \in \mathcal{N}(i)} k_\theta(x_i, x_j, e_{ij}) u_j^{(\ell)} w_{ij}.$$

- **Attention (nonlocal kernel)**: q, k, v maps with softmax over neighbors/modes (self- or cross-attn with a).
- **Hybrid/mixtures**: dynamic mixtures of fixed kernels, multigrid spectral+graph stacks, or alternating cross-attn refreshes from a .
- **Readout (projection to target channels)**

$$u(x) = Q(u^{(L)}(x), x), \quad Q : \mathbb{R}^C \rightarrow \mathbb{R}^{C_{\text{out}}} \text{ (pointwise)}$$

Objectives and Stability

Training objectives

- **Supervised:** $\|u_\theta - u_{\text{ref}}\|$ (e.g., L^1/L^2), spectral losses, gradient/energy terms.
- **Physics-informed:** PDE residuals $\|\mathcal{R}(u_\theta, a)\|$, BC penalties, integral constraints.

Stability & regularization

- *Residual/skip paths*, normalization (layer/instance), degree/softmax normalization on graphs.
- *Spectral control*: limit Fourier bandwidth (FNO), spectral norm/weight decay on \mathcal{W}^ℓ , anti-aliasing.
- *Lipschitz control*: clamp attention logits, bound dynamic mixing coefficients, Jacobian regularizers.

I/O variants

- *Bipartite input* \rightarrow *output*: first layer reads a at $\{\xi_j\} \rightarrow u$ at $\{x_i\}$, then propagate on u -graph.
- *Shared grid*: if $x \equiv \xi$, use a single graph/grid with concatenated channels for a .

FNO Parameter Complexity (channels C , layers L)

- **Setup.** Keep K_d low Fourier modes per axis $d = 1, \dots, D$. Number of learned spectral coefficients:

$$N_{\text{modes}} \approx \prod_{d=1}^D K_d \quad (\text{up to a small symmetry factor for real FFTs}).$$

- **Per Fourier layer** (width C) + Pointwise skip (1×1 conv):

$$\underbrace{W(k) \in \mathbb{C}^{C \times C}}_{\text{modewise multiplier}} + \text{Skip Connection} \quad \Rightarrow \quad \# \text{params} \approx 2 N_{\text{modes}} C^2 + C^2$$

- **Lift/Readout (once each):**

$$\#P = (C_{\text{in}} + \text{coord}) C \quad \#Q = C C_{\text{out}} \quad (\text{typically negligible vs. layers}).$$

Total (L layers):

$$\#\theta_{\text{FNO}} \approx L(2 N_{\text{modes}} + 1) C^2 + \#P + \#Q$$

Independent of grid size n (pixels/points).

Concrete Examples & Comparisons

- **Example A (2D):** $K_x = K_y = 12 \Rightarrow N_{\text{modes}} = 144$, $C = 64$, $L = 4$, $C_{\text{in}} = 2$, $C_{\text{out}} = 1$.

$$\#\text{per layer} \approx (2 \cdot 144 + 1) \cdot 64^2 = 289 \cdot 4096 \approx 1.18\text{M}$$

$$\#\text{total} \approx 4 \times 1.18\text{M} \approx 4.7\text{M} \quad (+ \#P + \#Q \approx \text{few} \times 10^2)$$

- **Example B (3D):** $K_x = K_y = K_z = 12 \Rightarrow N_{\text{modes}} = 1728$, $C = 64$, $L = 4$.

$$\#\text{per layer} \approx (2 \cdot 1728 + 1) \cdot 4096 \approx 14.15\text{M} \quad \Rightarrow \quad \#\text{total} \approx 56.6\text{M}.$$

- **Compute vs. params:**

- *Params:* depend on C , K_d , L (not grid).
- *FLOPs / pass:* FFT/IFFT $O(n \log n \cdot C)$ + low-mode multiplies $O(N_{\text{modes}} C^2)$.

- **Practical knobs**

- Reduce K_d (bandlimit) or C (width) to trade accuracy for size.
- Tie/share $W(k)$ across symmetric modes to cut constants $\sim 2\times$.
- Use bottleneck pointwise maps ($C \rightarrow C_b \rightarrow C$) to reduce C^2 .

Functional Tensor Decomposition

Siren Fourier Neural Operators (SirenFNOs): Re-design

- Consider the following Fourier layer

$$u^{(\ell+1)}(x) = \sigma \left(W^{(\ell)} u^{(\ell)}(x) + \mathcal{F}^{-1} \left(\mathcal{R}_\phi^{(\ell)}(k) \cdot \mathcal{F}(u^{(\ell)})(k) \right)(x) \right), \forall x \in \mathcal{D}$$

- In practical design, $\mathcal{R}_\phi^{(\ell)}$ is a complex tensor of size $K_1 \times K_2 \times \dots \times K_D \times C \times C$, where C means input and output channels
- SIREN employs sinusoidal activations to learn continuous implicit representations.
- A key advantage of SirenFNO is that the number of learnable parameters is fixed by the SIREN architecture and does not depend on the grid resolution.
- Thereby, our proposed SirenFNO efficiently represents both lower- and higher-frequency details across varying discretizations.

Siren Fourier Neural Operators (SirenFNOs): Re-design

- In our design, we take $\mathcal{R}_\phi^{(\ell)}$ as a parametric function $\Phi_\theta : [-1, 1]^D \rightarrow \mathbb{R}^{C \times C}$ with only number $|\theta|$ of parameters.
- Specifically, for each frequency mode \mathbf{k} ,

$$\Phi_\theta(\xi_k) = (\mathcal{O} \circ \phi^{(L)} \circ \dots \circ \phi^{(1)} \circ \mathcal{E}_B)(\xi_k),$$

where \mathcal{E}_B is a learnable random Fourier feature (RFF) embedding of the input frequency mode coordinates,

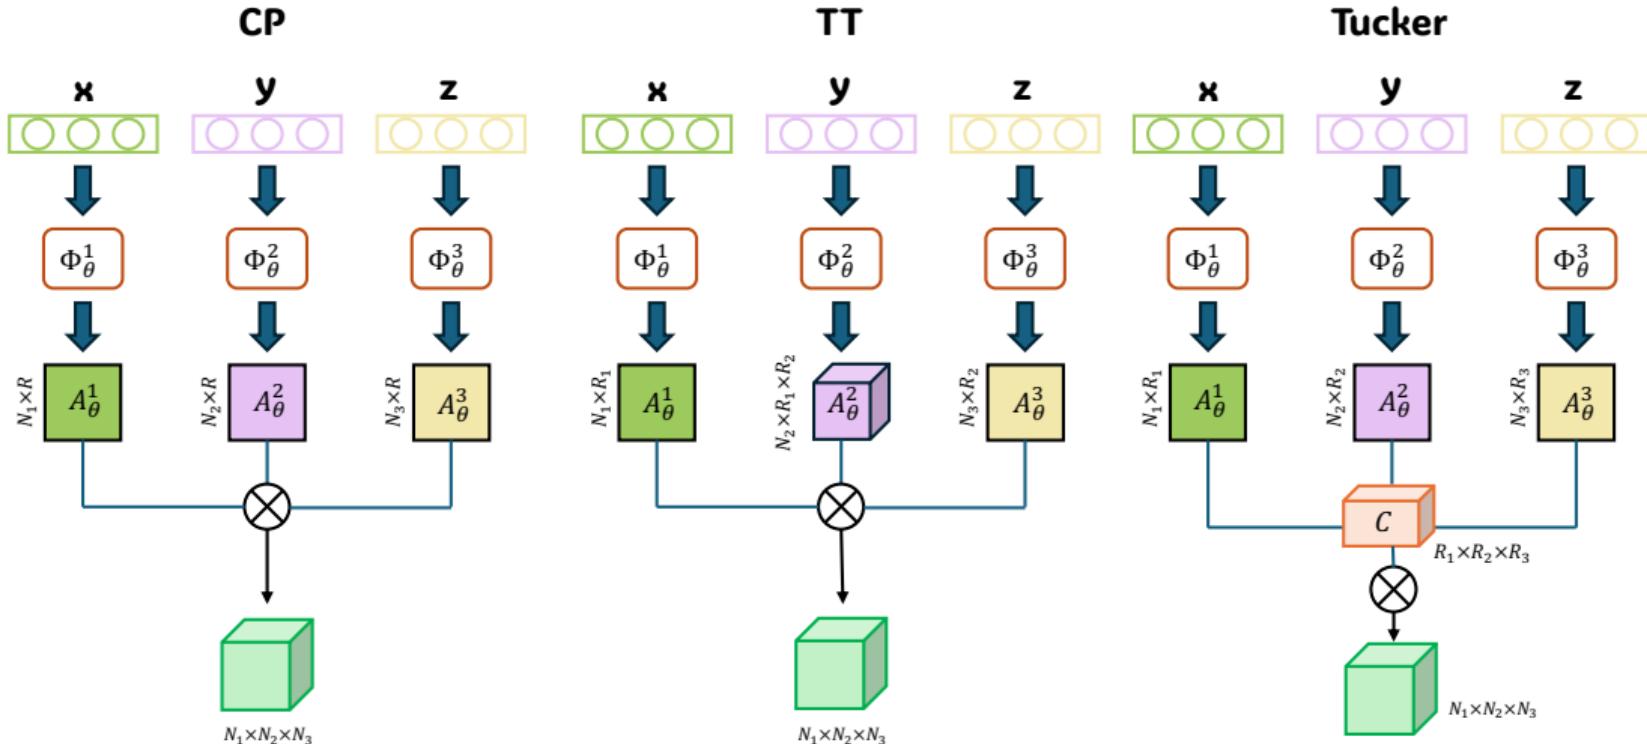
$$\mathcal{E}_B(\xi_k) = [\cos(\pi B^\top \xi_k), \sin(\pi B^\top \xi_k)] \in \mathbb{R}^{2m},$$

- $\phi^{(1)} : \mathbb{R}^{2m} \rightarrow \mathbb{R}^h$, and $\phi^{(\ell)} : \mathbb{R}^h \rightarrow \mathbb{R}^h$ for $\ell = 2, \dots, L$ each is a SIREN layer

$$\xi^{(\ell)} \leftarrow \phi^{(\ell)}(\xi_k^{(\ell-1)}) = \sin(w \cdot W^{(\ell)} \xi_k^{(\ell-1)} + b^{(\ell)}),$$

- And $\mathcal{O} : \mathbb{R}^h \rightarrow \mathbb{C}^{C \times C}$ is an MLP layer.

Functional Tensor Decomposition for SirenFNO



Benchmarks and Baselines

- 2D Darcy flow problem with Dirichlet boundary conditions

$$-\nabla \cdot (a(\mathbf{x}) \nabla u(\mathbf{x})) = f(\mathbf{x}), \mathbf{x} \in \mathcal{D}; \quad u(\mathbf{x}) = 0, \mathbf{x} \in \partial\mathcal{D}.$$

- Navier Stokes Equation in Velocity–pressure form (conceptual): Unknowns velocity $u(\mathbf{x}, t)$; given pressure $p(\mathbf{x}, t)$ and viscosity $\nu > 0$ and forcing $f(\mathbf{x}, t)$.

$$\begin{aligned}\partial_t u(\mathbf{x}, t) + (u \cdot \nabla) u(\mathbf{x}, t) - \nu \Delta u(\mathbf{x}, t) + \nabla p(\mathbf{x}, t) &= f(\mathbf{x}, t), \mathbf{x} \in \mathcal{D}, t > 0 \\ \nabla u(\mathbf{x}, t) &= 0, \mathbf{x} \in \mathcal{D}, t > 0; \quad u(\mathbf{x}, 0) = u_0(\mathbf{x})\end{aligned}$$

- 1-D Burgers' equation: **PDE** on $\mathcal{D} = (0, 1)$, $t \in [0, T]$ (periodic BCs):

$$\partial_t u(x, t) + u(x, t) \partial_x u(x, t) = \nu \partial_{xx} u(x, t), \quad u(x, 0) = u_0(x).$$

Name	Dataset	Resolution
Darcy	2-D Darcy flow	$32 \times 32, 128 \times 128$
NS	2-D Navier Stokes	128×128
Burgers	1-D Burgers' equation	$1024, T = 20$

Main Results: ℓ_2 Relative Test Errors

Dataset	Darcy		NS	Burgers	1DCFD	2DCFD	ReacDiff
Resolution	32 × 32	128 × 128	128 × 128	1024	1024	128 × 128	1024
FNO	7.30e-2	6.26e-2	5.61e-2	1.12e-2	8.27e-3	1.64e-2	1.18e-4
UFNO	7.07e-2	6.00e-2	5.37e-2	1.70e-2	8.29e-3	1.21e-2	1.98e-4
U-NO	6.57e-2	7.68e-2	4.96e-2	7.53e-3	7.62e-3	1.62e-2	7.79e-5
AM-FNO (MLP)	3.72e-2	5.10e-2	4.02e-2	8.92e-3	7.89e-3	1.21e-2	1.78e-3
SirenFNO [†]	6.32e-2	5.82e-2	5.52e-2	8.21e-3	7.41e-3	1.15e-2	7.40e-5
CP-SirenFNO [†]	4.99e-2	3.96e-2	5.04e-2	8.12e-3	5.91e-3	7.28e-3	8.98e-5
TT-SirenFNO [†]	5.42e-2	5.36e-2	5.87e-2	7.94e-3	5.58e-3	8.75e-3	8.51e-5
Tucker-SirenFNO [†]	4.45e-2	4.51e-2	4.19e-2	7.18e-3	7.31e-3	7.30e-3	7.81e-5
SirenFNO	3.51e-2	2.36e-2	3.51e-2	5.52e-3	6.77e-3	6.83e-3	6.31e-5
CP-SirenFNO	4.04e-2	3.03e-2	3.78e-2	8.48e-3	4.75e-3	7.40e-3	7.58e-5
TT-SirenFNO	4.18e-2	3.33e-2	3.24e-2	7.52e-3	6.32e-3	7.41e-3	8.92e-5
Tucker-SirenFNO	3.98e-2	3.18e-2	4.67e-2	8.11e-3	5.29e-3	6.90e-3	8.71e-5

ℓ_2 relative test errors. All models are trained and tested on the same corresponding resolution. Models highlighted with [†] are using the identical architecture with FNO for ablation study. SirenFNO and its variants in boldface but without [†] are our proposed methods.

Main Results: Number of Parameters

Dataset	Darcy		NS	Burgers	1DCFD	2DCFD	ReacDiff
Resolution	32 × 32	128 × 128	128 × 128	1024	1024	128 × 128	1024
FNO	1192.8	1192.8	4469.6	4216.2	4216.2	4469.9	4216.2
UFNO	5876.2	5876.2	990.8	4258.5	1892.2	991.0	1892.2
U-NO	1792.7	4250.3	8726.8	1726.1	1726.1	7602.3	5658.2
AM-FNO (MLP)	385.5	385.5	4443.1	207.6	823.1	385.6	823.1
SirenFNO	308.9	308.9	579.5	308.9	304.8	304.8	304.8
CP-SirenFNO	63.9	63.9	138.9	70.1	57.7	64.0	57.7
TT-SirenFNO	92.5	109.2	211.6	84.5	72.0	92.7	72.0
Tucker-SirenFNO	162.6	162.6	596.4	74.2	61.8	96.8	61.8

Time-aware Neural Operators

Example PDE: 1D Viscous Burgers

- **PDE** on $\mathcal{D} = (0, 1)$, $t \in [0, T]$ (periodic BCs):

$$\partial_t u(x, t) + u(x, t) \partial_x u(x, t) = \nu \partial_{xx} u(x, t), \quad u(x, 0) = u_0(x).$$

- **Task (operator learning).** Learn $\mathcal{G}_\theta : (u_0, \nu) \mapsto u(\cdot, t)$

either $\mathcal{G}_\theta^{\text{traj}} : (u_0, \nu) \mapsto u(\cdot, \cdot)$ or $\mathcal{G}_\theta^{\Delta t} : (u(\cdot, t), \nu) \mapsto u(\cdot, t + \Delta t)$.

- **Data (typical).** Draw u_0 from a random band-limited field; simulate with a spectral solver (forward Euler/ETDRK4) at high resolution; subsample in space/time for training.
- **Loss.** Supervised spacetime error L_2 (optionally + physics residuals/BC penalties).

Time-Aware Neural Operators: Three Blueprints

- **Spacetime FNO (t as another axis)**

$$\boxed{\text{Lift: }} u^{(0)}(x, t) = P[u_0(x), x, t, \nu]$$

$$\boxed{\text{Fourier layer: }} \widehat{v}^{(\ell)}(k_x, k_t) = W_\ell(k_x, k_t) \widehat{u}^{(\ell)}(k_x, k_t) \text{ for low } (k_x, k_t)$$

$$u^{(\ell+1)} = \sigma(\mathcal{F}^{-1}[\widehat{v}^{(\ell)}] + Bu^{(\ell)}), \quad u = Q u^{(L)}.$$

Pros: global coupling in space *and* time; single forward pass for $u(x, t)$. *Notes:* keep few k_t modes; window/pad for non-periodic t .

- **Autoregressive (one-step) Operator**

$$\phi_{\Delta t} : u(\cdot, t) \mapsto u(\cdot, t+\Delta t) \quad (\text{FNO/GNO/Attn backbone})$$

$$u_{n+1} = \phi_{\Delta t}(u_n; \nu), \quad n = 0, \dots, N-1.$$

Pros: flexible horizons; plug-in with variable Δt . *Training:* teacher forcing & scheduled sampling; rollout loss across multiple steps.

Time-Aware Neural Operators: Three Blueprints

- **Graph/Attention Neural Operator with Time** *Pros:* irregular meshes, nonuniform time grids, complex BC/geometry.

- **Spacetime discretization.** Let $X = \{x_i\}_{i=1}^M$, $T = \{t_n\}_{n=0}^N$. Create nodes

$\mathcal{V} = \{(i, n) : x_i \in X, t_n \in T\}$ and edges \mathcal{E} by, *past light-cone and enforcing causality with $t_m \leq t_n$,*

$$\mathcal{N}(i, n) = \{(j, m) : \|x_i - x_j\| \leq r_x, 0 \leq t_n - t_m \leq r_t\}$$

- **Node features.** $h_{i,0}^{(0)} = P(u_0(x_i), x_i, t_0, \nu)$; for $n > 0$, initialize $h_{i,n}^{(0)} = P(x_i, t_n, \nu)$.
- **Separable space–time operator (efficient).**

$$h_{i,n}^{(\ell+1)} = \sigma \left(\mathcal{W}^{(\ell)} h_{i,n}^{(\ell)} + \sum_{j \in \mathcal{N}_x(i)} \kappa_s^{(\ell)}(\phi_{ij}; h_{i,n}^{(\ell)}, h_{j,n}^{(\ell)}) U_s^{(\ell)} h_{j,n}^{(\ell)} w_{ij}^x + \sum_{m \in \mathcal{N}_t(n)} \kappa_t^{(\ell)}(\tau_{nm}; h_{i,n}^{(\ell)}, h_{i,m}^{(\ell)}) U_t^{(\ell)} h_{i,m}^{(\ell)} w_{nm}^t \right)$$

with $\mathcal{N}_t(n) \subseteq \{m : 0 \leq t_n - t_m \leq r_t\}$ and (optional) cross term:

$$+ \sum_{j,m} \beta^{(\ell)}(\phi_{ij}, \tau_{nm}) U_c^{(\ell)} h_{j,m}^{(\ell)}$$

Burgers with FNO: Two Concrete Setups

① Spacetime FNO (2D FFT over (x, t)) (showing one layer)

$$\begin{aligned}\widehat{v}(k_x, k_t) &= W(k_x, k_t) \widehat{u}(k_x, k_t) \quad \text{for } |k_x| \leq K_x, \quad |k_t| \leq K_t. \\ u' &= \sigma(\mathcal{F}^{-1}[\widehat{v}] + Bu),\end{aligned}$$

BCs in t : pad/window or cosine transform along t if non-periodic.

② Autoregressive Spatial-FNO (1D FFT over x)

$$\boxed{\text{Lift}} : \quad u^{(0)}(x) = P(u_n(x), x, \nu, \Delta t) \in \mathbb{R}^C$$

For $\ell = 0, \dots, L-1$:

$$\widehat{v}^{(\ell)}(k_x) = W_\ell(k_x) \widehat{u}^{(\ell)}(k_x) \quad \text{for } |k_x| \leq K_x, \quad u^{(\ell+1)}(x) = \sigma(\mathcal{F}_x^{-1}[\widehat{v}^{(\ell)}] + B_\ell u^{(\ell)}(x))$$

$$\boxed{\text{Readout}} : \quad \tilde{u}_{n+1}(x) = Q u^{(L)}(x), \quad \boxed{\text{Stabilize}} : \quad u_{n+1} = u_n + \gamma \tilde{u}_{n+1} \quad (\text{residual gain } \gamma \in (0, 1]).$$

- *Stable rollouts*: spectral padding, limit low modes, small residual gain.
- *Curriculum*: train 1-step, then 4-step, then full horizon.

Training, Physics, and Practical Tips (Burgers)

- Predict full trajectory $u(\cdot, t)$ or endpoints $u(\cdot, T)$.

$$\mathcal{L} = \underbrace{\|u_\theta - u_{\text{ref}}\|_{L^2}}_{\text{data}} + \lambda_{\text{PDE}} \underbrace{\|\partial_t u_\theta + u_\theta \partial_x u_\theta - \nu \partial_{xx} u_\theta\|_{L^2}}_{\text{physics}} + \lambda_{\text{BC}} \|\text{BC}(u_\theta)\|.$$

- or rollout

$$\mathcal{L} = \underbrace{\sum_{n=0}^{N-1} \alpha_n \|u_{n+1} - \phi_{\Delta t}(u_n)\|_{L^2}^2}_{\text{data rollout}} + \lambda_{\text{PDE}} \underbrace{\|\partial_t u_\theta + u_\theta \partial_x u_\theta - \nu \partial_{xx} u_\theta\|_{L^2}}_{\text{physics}} + \lambda_{\text{BC}} \|\text{BC}(u_\theta)\|.$$

Training, Physics, and Practical Tips

Stability & generalization

- Limit temporal bandwidth (K_t) or use AR with small Δt .
- Anti-aliasing / spectral filtering; weight sharing across time blocks.
- Randomize grids (space/time) at train time for resolution-agnosticity.

When to pick which

- *Smooth global dynamics, fixed T :* spacetime FNO (fast, single pass).
- *Variable horizons or control-in-time:* autoregressive FNO/GNO.
- *Irregular meshes or moving boundaries:* graph/attention operator with temporal links.

Neural Operators for SDE/SPDEs

From ODE/PDEs to SDE/SPDEs: What Changes?

- So far: **deterministic PDEs**, e.g.,

$$\partial_t u = \mathcal{F}(u, a), \quad u(\cdot, 0) = u_0.$$

Neural operator learns

$$\mathcal{G}_\theta : (u_0, a) \mapsto u(\cdot, t).$$

- For **stochastic differential equations (SDEs)**, dynamics are driven by noise:

$$dX_t = f(X_t, t) dt + g(X_t, t) dW_t,$$

where W_t is Brownian motion (or another stochastic process).

- **Stochastic PDE (SPDE):** add noise in time and space, e.g.

$$\partial_t u(x, t) = \mathcal{F}(u, a)(x, t) + \mathcal{G}(u, a)(x, t) \dot{W}(x, t),$$

where $\dot{W}(x, t)$ denotes space-time noise (formally the derivative of a space-time Wiener process, often idealised as space-time white noise).

From ODE/PDEs to SDE/SPDEs: What Changes?

- The solution is a **random path** $t \mapsto X_t(\omega)$. We care about:
 - individual sample paths (pathwise behavior),
 - or statistics: $\mathbb{E}[X_t]$, $\text{Var}[X_t]$, full law p_{X_t} .
- Neural operators can learn the **solution map** (Itô map):

$$\mathcal{G} : (X_0, \text{noise}) \mapsto X(\cdot),$$

using the same operator backbones (FNO, DeepONet, GNO) we use for PDEs.

- Solution is now a **random field** $(x, t) \mapsto u(x, t, \omega)$. We can learn:
 - *pathwise map*: a specific noise realization \Rightarrow a field,
 - or *statistical map*: $(x, t) \mapsto$ quantities like $\mathbb{E}[u(x, t)]$, $\text{Var}[u(x, t)]$, or even the full distribution of $u(x, t)$
- Neural operators still target an operator

$$\mathcal{G} : (u_0, \text{forcing/noise}) \mapsto u(\cdot, \cdot),$$

just now in *space-time*.

SDE Solution as an Operator (Pathwise View)

- Consider an Itô SDE on $[0, T]$:

$$dX_t = f(X_t, t; \theta) dt + g(X_t, t; \theta) dW_t, \quad X_0 \sim p_0.$$

- After discretization (e.g. Euler–Maruyama with step Δt):

$$X_{n+1} = X_n + f(X_n, t_n) \Delta t + g(X_n, t_n) \Delta W_n,$$

where $\{\Delta W_n\}$ are Gaussian increments.

- Pathwise solution map (Itô map):**

$$\mathcal{G} : (X_0, \{\Delta W_n\}_{n=0}^{N-1}) \longmapsto \{X_n\}_{n=0}^N.$$

Deterministic map once the noise path is fixed.

- Neural operator perspective:**

- Input function(s): noise path, control signals, parameters θ .
- Output function: state trajectory $X(\cdot)$ (possibly multi-dimensional).
- Use a time-aware operator backbone (spacetime FNO, AR-FNO, GNO-in-time) to approximate \mathcal{G} .

Neural Operator Blueprints for SDEs

① Pathwise SDE Neural Operator (trajectory generator)

$$\mathcal{G}_\theta : (X_0, \{\Delta W_n\}) \mapsto \{X_n\}_{n=0}^N$$

- Treat $\{\Delta W_n\}$ as an *input function in time*.
- Use spacetime FNO or autoregressive FNO/GNO as in the Burgers setup.
- Sample new trajectories by sampling new noise paths.

② Statistic / moment operator

$$\mathcal{H}_\theta : (X_0, \theta) \mapsto (m(t), s^2(t))_{t \in [0, T]}$$

- Learn map to mean/variance curves or other summary functionals.
- Useful when full path samples are not required.

③ Example: OU process with FNO

$$dX_t = \alpha(\mu - X_t) dt + \sigma dW_t.$$

SPDE Solution as a Space–Time Operator

- Consider an SPDE on $\mathcal{D} \times [0, T]$:

$$\partial_t u(x, t) = \mathcal{L}u(x, t) + \mathcal{N}(u, a)(x, t) + \mathcal{G}(u, a)(x, t) \dot{W}(x, t),$$

with $u(x, 0) = u_0(x)$ and boundary conditions on $\partial\mathcal{D}$.

- After discretization (space + time), e.g. finite difference + Euler–Maruyama:

$$u_{n+1} = \Phi(u_n, a, \xi_n),$$

where $\xi_n(x)$ are random noise fields on the spatial grid at time t_n .

- Pathwise solution map** (Itô map for SPDE):

$$\mathcal{G} : (u_0, \{\xi_n\}_{n=0}^{N-1}) \longmapsto \{u_n\}_{n=0}^N.$$

Once the noise fields $\{\xi_n\}$ are fixed, this is deterministic.

- Neural operator perspective:**

- Inputs are functions on space–time: $u_0(x)$, $a(x)$, noise field(s).
- Output is a function on space–time: $u(x, t)$.
- Use the same backbones: FNO, DeepONet, GNO, attention operators.

Neural Operator Blueprints for SPDEs (FNO / GNO)

① Spacetime FNO with stochastic forcing

- Treat (x, t) as a 2D domain and stack channels:

channels: $[u(x, t), a(x), \text{noise}(x, t), \text{coords}]$.

- Apply 2D FNO layers over (x, t) :

$$u^{(\ell+1)} = \sigma\left(W_0 u^{(\ell)} + \mathcal{F}^{-1}\left[W_\ell(k_x, k_t) \widehat{u^{(\ell)}}\right]\right).$$

- Learn the map $(u_0, a, \text{noise}) \mapsto u(\cdot, \cdot)$ in one shot.

② Autoregressive spatial operator with noisy increments

- Time is stepped; space is handled by FNO / GNO:

$$u_{n+1}(x) = \phi_\theta(u_n(\cdot), a(\cdot), \xi_n(\cdot))(x),$$

where $\xi_n(\cdot)$ is the noise field at step n .

- Roll out $u_0 \rightarrow u_1 \rightarrow \dots \rightarrow u_N$; train with multi-step losses.

SDE vs SPDE Neural Operators (What's Really Different?)

- **SDE:** state is finite-dimensional, $X_t \in \mathbb{R}^d$.

$$dX_t = f(X_t, t) dt + g(X_t, t) dW_t.$$

$$\mathcal{G}_\theta : (X_0, \{\Delta W_n\}) \mapsto \{X_n\}_{n=0}^N.$$

- Operator on *time paths* $[0, T] \rightarrow \mathbb{R}^d$.
- Use 1D (in t) operator backbones (FNO-in-time, AR operators).
- **SPDE:** state is a field, $u : \mathcal{D} \times [0, T] \rightarrow \mathbb{R}^m$.

$$\partial_t u(x, t) = \mathcal{F}(u, a)(x, t) + \mathcal{G}(u, a)(x, t) \dot{W}(x, t).$$

$$\mathcal{G}_\theta : (u_0(\cdot), a(\cdot), \{\xi_n(\cdot)\}) \mapsto \{u_n(\cdot)\}_{n=0}^N.$$

- Operator on *space-time fields* $\mathcal{D} \times [0, T] \rightarrow \mathbb{R}^m$.
- Use spacetime FNO, or spatial FNO/GNO + temporal stepping.
- **Same philosophy:** learn the Itô map (noise \Rightarrow solution), but on different function spaces (time vs space-time).

Q & A

All questions are welcome. Thank you!